


From Scramble to Solution

Developing an Optimal Solver for
the Kilominx, a Rubik's Cube Variant

Samuel Borg

Supervisor: Ian Gent

March 26, 2025



School of Computer Science
University of St Andrews

Abstract

The Rubik's Cube is the most well-known combination puzzle of all time, and has inspired numerous variants of the puzzle. One such puzzle is the Kilominx, a 2x2 dodecahedron-shaped puzzle with 12 faces and 20 cubies (individual cube elements of the puzzle). While the Rubik's Cube has already been extensively researched, the Kilominx has not.

This report details my work in creating an optimal solver for the Kilominx, guaranteeing solutions in as few moves as possible. A separate program I wrote precomputes several pattern databases, each containing the number of moves required to solve different sets of cubies in all possible states. The solver uses an iterative-deepening A* (IDA*) algorithm to find the optimal solution, using the pattern databases as a lower-bound heuristic. In order to efficiently index into the pattern databases, the solver computes a Lehmer code for the puzzle state, which is then used to calculate the state's sequential index.

In theory, the solver can optimally solve any state given enough time, but in practice it can only find solutions up to a depth of 14 within a reasonable timeframe. As solutions require longer sequences of moves to solve, the solver takes exponentially more time to find the optimal solution. With further optimisation, the solver could be used to help tighten the bounds on the Kilominx's God's Number - the maximum number of moves needed to solve the puzzle from any given state.

Declaration

I declare that the material submitted for assessment is my own work except where credit is explicitly given to others by citation or acknowledgement. This work was performed during the current academic year except where otherwise stated.

The main text of this project report is [insert word count] words long, including project specification and plan.

In submitting this project report to the University of St Andrews, I give permission for it to be made available for use in accordance with the regulations of the University Library. I also give permission for the title and abstract to be published and for copies of the report to be made and supplied at cost to any bona fide library or research worker, and to be made available on the World Wide Web. I retain the copyright in this work.

Acknowledgements

I would like to thank my supervisor, Ian Gent, for his invaluable guidance, insight and support throughout this project. His expertise and enthusiasm have made this project an incredibly engaging and rewarding journey.

I would also like to thank my parents, Juliet and Peter, my brother, Tom, and my amazing friends Emil, Ben, Zoe, Josh, Maya, Jess and Alfie for their unwavering support and belief in me all the way through my studies.

Contents

1	Introduction	5
1.1	Objectives	5
2	Context Survey	6
2.1	Korf's Algorithm for the Rubik's Cube	6
2.1.1	Iterative-Deepening A*	6
2.1.2	Pattern Databases	6
2.2	Sequential Indexing with Lehmer Codes	6
2.3	God's Number	6
3	Requirements Specification	7
4	Design	8
5	Implementation	9
6	Evaluation	10
6.1	Experimental Results	10
6.2	Evaluation Against Objectives	10
6.3	Critical Appraisal	10
7	Conclusions	11
A	Testing Summary	12
B	User Manual	13
	References	14

Chapter 1

Introduction

When the Rubik's Cube was globally released in 1980, it became an instant success, selling around 200 million units by the end of 1983 [1]. The Rubik's Cube also quickly became a popular subject of research for computer scientists, in part due to the massive 43 quintillion possible states [2] that the cube can be in. Many different algorithms were designed to try and solve the cube in as few a number of moves as possible, but it wasn't until 1997 that Richard E. Korf published a paper [3] describing a method to solve the Rubik's Cube optimally (the shortest possible number of moves to solve any given cube state) by using large lookup tables called pattern databases[4] as a heuristic function to guide an IDA* search algorithm.

The widespread success of the Rubik's Cube also led to the creation of a number of different variants of the puzzle, each working similarly to the Rubik's Cube but of different shapes and sizes. The Kilominx is one of these variants, part of the larger "minx" family of dodecahedron-shaped puzzles. The Kilominx is a 2x2 dodecahedral puzzle with 12 faces and 20 cubies (the individual cube-shaped elements of the puzzle). Although some research has been done on tightening the bounds of its God's Number, no optimal solver has previously been created for the puzzle.



Figure 1: An image of a Kilominx.

1.1 Objectives

Chapter 2

Context Survey

2.1 Korf's Algorithm for the Rubik's Cube

2.1.1 Iterative-Deepening A*

2.1.2 Pattern Databases

2.2 Sequential Indexing with Lehmer Codes

2.3 God's Number

Chapter 3

Requirements Specification

Chapter 4

Design

Chapter 5

Implementation

Chapter 6

Evaluation

6.1 Experimental Results

6.2 Evaluation Against Objectives

6.3 Critical Appraisal

Chapter 7

Conclusions

Appendix A

Testing Summary

Appendix B

User Manual

References

- [1] *The Lighter Side of Mathematics*. URL: <https://archive.org/details/lightersideofmat0000unse/page/340/mode/1up>. Last accessed on 25/03/25.
- [2] *Mathematics of the Rubik's Cube*. URL: <https://ruwix.com/the-rubiks-cube/mathematics-of-the-rubiks-cube-permutation-group/>. Last accessed on 25/03/25.
- [3] R. E. Korf. “Finding Optimal Solutions to Rubik’s Cube Using Pattern Databases”. In: *AAAI’97*. 1997. URL: <https://www.cs.princeton.edu/courses/archive/fall06/cos402/papers/korfrubik.pdf>. Last accessed on 25/03/25.
- [4] J. C. Culberson and J. Schaeffer. “Pattern Databases”. In: *Computational Intelligence* 14.3 (1998). URL: <https://onlinelibrary.wiley.com/doi/abs/10.1111/0824-7935.00065>. Last accessed on 26/03/25.